

Peter Van Sandt

✉ peter.vansandt@gmail.com
📁 [eyepatchParrot.github.io](https://github.com/eyepatchParrot)
in [linkedin.com/in/petervansandt](https://www.linkedin.com/in/petervansandt)

WORK EXPERIENCE

Oct 2019 - **Dell-EMC Isilon**, *Senior Software Engineer*, Seattle, WA.

Present

Jun 2018 - *Software Engineer - Performance.*

Oct 2019

2017 *Software Engineer Intern - Performance.*

- Designed and developed request handling, parsing, and dispatch in development of forthcoming protocol stack: minimizing data copies and branches in create, read, delete, and list operations.
- Developed performance test plan and corresponding benchmarks to provide performance comparison and analysis during development of forthcoming protocol stack: enabling the team to match the performance of an industry-leading NFS implementation.
- Researched and regularly presented performance and architectural analysis of academic and competing stacks to leadership team: focusing on concurrent creates in parallel file systems and throughput in object storage.
- Developed and extended non-intrusive, trace-based and low-overhead, distribution-based workload analysis tools with dtrace and SQL, so that customers could size their clusters 200x faster and support engineers could non-disruptively assess on-going workload distributions.
- Developed and extended automated performance analysis tooling with pandas: modeling read path on a distributed file system to highlight reads that use caches inefficiently, and identifying a 50% performance regression.
- Root cause analysis of performance issues in an NFS protocol stack and authentication reducing CPU utilization in hot spots and improve worst case performance.

Oct 2016 - **University of Wisconsin-Madison, Computer Sciences**, *Undergraduate Researcher.*

May 2018

- Developed interpolation search algorithms for uniform and non-uniform data to use constant factor improvements and asymptotic improvements using the distribution of the data to improve throughput in numpy by 380% and LevelDB by 146%.
- Optimized LevelDB by replacing use of mutexes in the fast-path for reads with copy-on-write version updates and lock-free LRU improving read throughput by up to 45%.

Jun 2016 - **Jump Trading, LLC.**, *Software Developer Intern*, Chicago, IL.

Aug 2016

- Eliminated algorithmic bottlenecks in 100KLOC internal application improving response time by 200%.
- Developed performance visualization and analysis tool for internal use.

EDUCATION

May 2018 **B.S., Computer Sciences, Honors, ΦBK**, *University of Wisconsin-Madison*, GPA: 4.0/4.0.

- Graduate Coursework: Advanced Algorithms, High Performance Computing, Distributed Systems
- Languages: C++, C, Python 2.7, dtrace, AVX2, CUDA C

AWARDS

2019 Dell Global Storage Hackathon, Best in *AI/ML* and *Innovative Idea*

2019 Van Sandt, Peter, Yannis Chronis, and Jignesh M. Patel. "Efficiently Searching In-Memory Sorted Arrays: Revenge of the Interpolation Search?." *Proceedings of the 2019 International Conference on Management of Data*. 2019.

2017 DeWitt Undergraduate Scholarship for academic excellence and research in Computer Science

PERSONAL PROJECTS

2017 Developed B-tree index, buffer manager, and space efficient word search for database class covering database internals, design, and algorithms.

2016 Vectorized insertion and rank order sorts of fixed-size arrays using C and AVX2 intrinsics.

2016 Developed text-record sorting utility using trie-based burst-sort in C winning as the fastest in a class of 300.

2012 Developed GPU-accelerated, Toom-Cook multiplication bignum library in C++, and CUDA C.